


In [1]:

```
#Loading libraries and data
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
df=pd.read_csv('/home/sandhan/Downloads/WA_Fn-UseC_-Telco-Customer-Churn.csv') #Loading data
df
```

Out[1]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleL
0	7590-VHVEG	Female	0	Yes	No	1	No	No pl se
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	No pl se
4	9237-HQITU	Female	0	No	No	2	Yes	
...	
7038	6840-RESVB	Male	0	Yes	Yes	24	Yes	
7039	2234-XADUH	Female	0	Yes	Yes	72	Yes	
7040	4801-JAZL	Female	0	Yes	Yes	11	No	No pl se
7041	8361-LTMKD	Male	1	Yes	No	4	Yes	
7042	3186-AJIEK	Male	0	No	No	66	Yes	

7043 rows × 21 columns



In [2]:

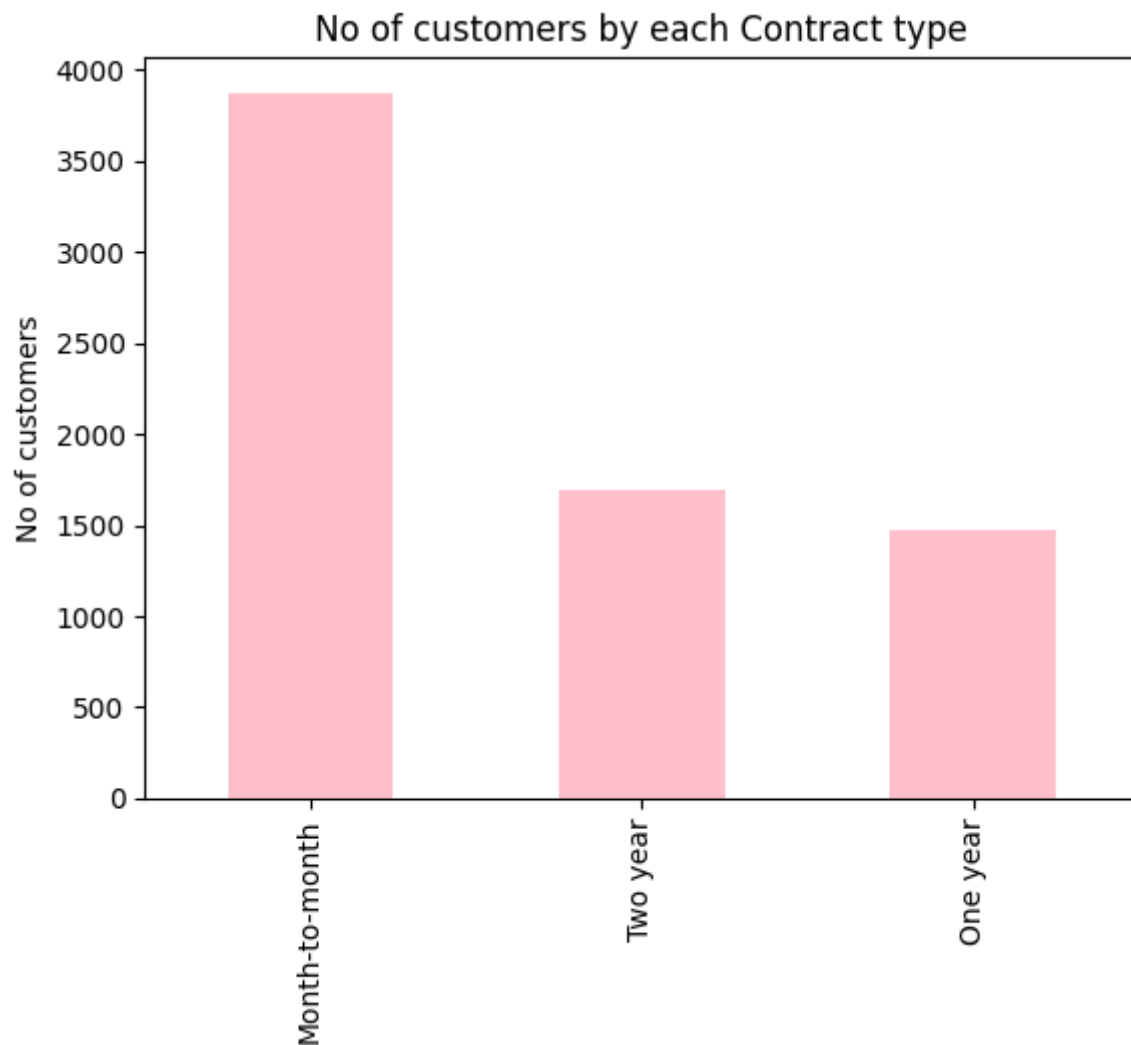
```
#DATA VISUALISATION
```

In [3]:

```
#Number of customers with each contract type  
plot1=df["Contract"].value_counts().plot(kind='bar',color='pink')  
plot1.set_ylabel("No of customers")  
plot1.set_title("No of customers by each Contract type")
```

Out[3]:

Text(0.5, 1.0, 'No of customers by each Contract type')




In [4]:

```
#Number of customers with each InternetService Type  
plot2=df["InternetService"].value_counts().plot(kind='bar',color='grey')  
plot2.set_ylabel("No of customers")  
plot2.set_title("No of customers with each InternetService Type")
```

Out[4]:

Text(0.5, 1.0, 'No of customers with each InternetService Type')



In [5]:

```
#No of customers by their tenure
plot3=sns.distplot(df['tenure'], hist=True, kde=False, bins=int(180/5),\
                    color='darkred', hist_kws={'edgecolor': 'black'})
plot3.set_ylabel('No of Customers')
plot3.set_xlabel('Tenure (months)')
plot3.set_title('No of Customers by their tenure')
```

/tmp/ipykernel_7786/2452378685.py:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

```
plot3=sns.distplot(df['tenure'], hist=True, kde=False, bins=int(180/5),\
```

Out[5]:

Text(0.5, 1.0, 'No of Customers by their tenure')




In [6]:

#Association between TotalCharges and MonthlyCharges

```
df['TotalCharges']=df['TotalCharges'].str.replace(' ','0')  
df['TotalCharges']=df['TotalCharges'].str.replace('0 0','0')  
df['TotalCharges']=df['TotalCharges'].astype(float)  
df['TotalCharges']=df['TotalCharges'].replace('0',df['TotalCharges'].mean())  
sns.regplot(x=df['MonthlyCharges'],y=df['TotalCharges'],color='orange')
```

Out[6]:

<AxesSubplot: xlabel='MonthlyCharges', ylabel='TotalCharges'>

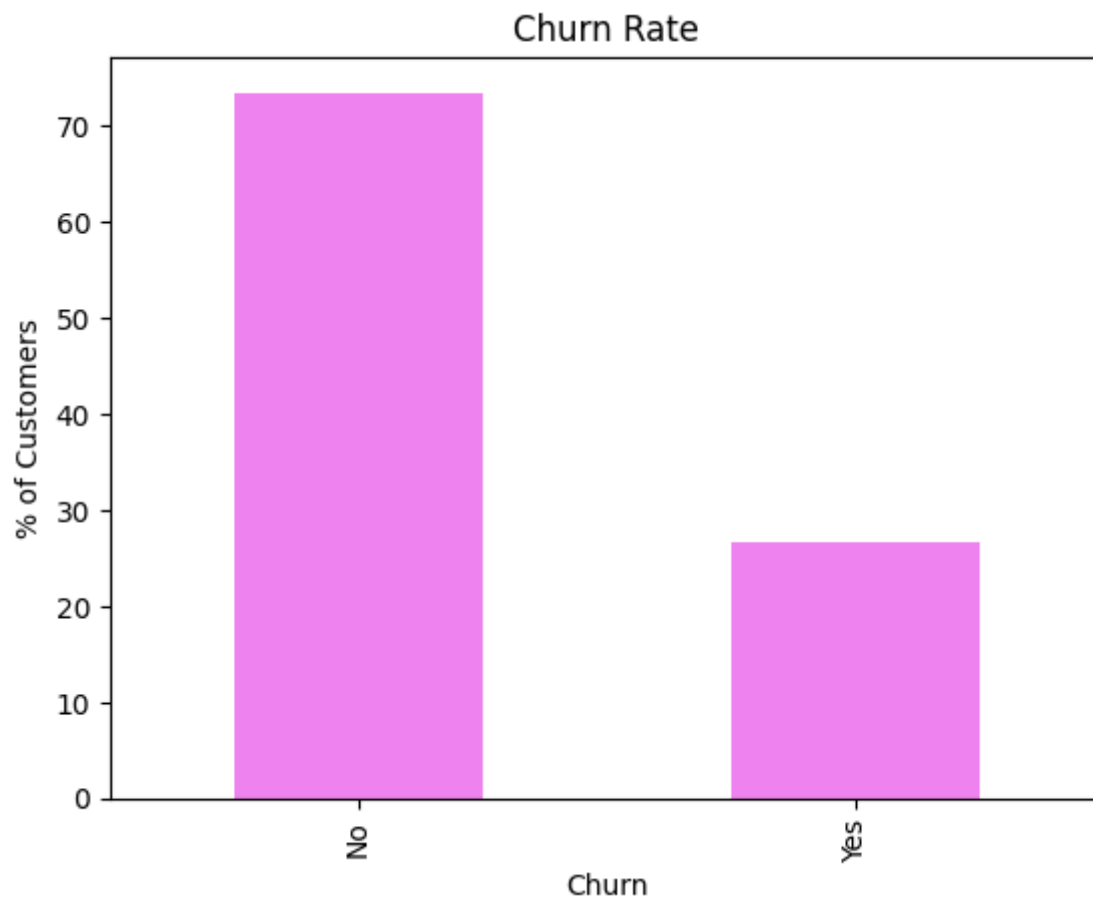


In [7]:

```
#Churn rate of customers with percentage  
plot4=(df['Churn'].value_counts()*100/len(df)).plot(kind='bar',color='violet')  
plot4.set_ylabel('% of Customers')  
plot4.set_xlabel('Churn')  
plot4.set_title('Churn Rate')
```

Out[7]:

Text(0.5, 1.0, 'Churn Rate')




In [8]:

```
#Churn rate of customers with each contract type
```

```
contract_churn=df.groupby(['Contract','Churn']).size().unstack()
```

```
plot6=(contract_churn.T*100.0 / contract_churn.T.sum()).T.plot(kind='bar',width=0.3
```



In [9]:

```
#UNDERSTANDING DATA
```

In [10]:

```
df.shape
```

Out[10]:

(7043, 21)

In [11]:

```
df.columns
```

Out[11]:

```
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependent  
s',  
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',  
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupp  
ort',  
      'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBillin  
g',  
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],  
      dtype='object')
```

In [12]:

```
df.dtypes
```

Out[12]:

```
customerID      object  
gender          object  
SeniorCitizen   int64  
Partner         object  
Dependents      object  
tenure          int64  
PhoneService    object  
MultipleLines   object  
InternetService object  
OnlineSecurity  object  
OnlineBackup    object  
DeviceProtection object  
TechSupport     object  
StreamingTV     object  
StreamingMovies object  
Contract        object  
PaperlessBilling object  
PaymentMethod   object  
MonthlyCharges  float64  
TotalCharges    float64  
Churn           object  
dtype: object
```

In [13]:

```
df.drop(['customerID'],axis=1,inplace=True)
```


In [14]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   gender                7043 non-null  object 
 1   SeniorCitizen         7043 non-null  int64  
 2   Partner               7043 non-null  object 
 3   Dependents            7043 non-null  object 
 4   tenure                7043 non-null  int64  
 5   PhoneService          7043 non-null  object 
 6   MultipleLines         7043 non-null  object 
 7   InternetService       7043 non-null  object 
 8   OnlineSecurity        7043 non-null  object 
 9   OnlineBackup          7043 non-null  object 
10   DeviceProtection     7043 non-null  object 
11   TechSupport           7043 non-null  object 
12   StreamingTV           7043 non-null  object 
13   StreamingMovies       7043 non-null  object 
14   Contract              7043 non-null  object 
15   PaperlessBilling      7043 non-null  object 
16   PaymentMethod         7043 non-null  object 
17   MonthlyCharges        7043 non-null  float64 
18   TotalCharges          7043 non-null  float64 
19   Churn                 7043 non-null  object 
dtypes: float64(2), int64(2), object(16)
memory usage: 1.1+ MB
```

In [15]:

```
#DATA MANIPULATION
```


In [16]:

```
#Label Encoding
from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
cols=['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines', 'InternetService']
for i in cols:
    df[i]=encoder.fit_transform(df[i])
df
```

Out[16]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	Internet
0	0	0	1	0	1	0	1	
1	1	0	0	0	34	1	0	
2	1	0	0	0	2	1	0	
3	1	0	0	0	45	0	1	
4	0	0	0	0	2	1	0	
...	
7038	1	0	1	1	24	1	2	
7039	0	0	1	1	72	1	2	
7040	0	0	1	1	11	0	1	
7041	1	1	1	0	4	1	2	
7042	1	0	0	0	66	1	0	

7043 rows × 20 columns



In [17]:

```
df.dtypes
```

Out[17]:

```
gender                int64
SeniorCitizen         int64
Partner               int64
Dependents             int64
tenure                int64
PhoneService          int64
MultipleLines         int64
InternetService       int64
OnlineSecurity        int64
OnlineBackup          int64
DeviceProtection      int64
TechSupport           int64
StreamingTV           int64
StreamingMovies       int64
Contract              int64
PaperlessBilling       int64
PaymentMethod         int64
MonthlyCharges        float64
TotalCharges          int64
Churn                 int64
dtype: object
```

In [18]:

```
df.describe()
```

Out[18]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Multi
count	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043
mean	0.504756	0.162147	0.483033	0.299588	32.371149	0.903166	C
std	0.500013	0.368612	0.499748	0.458110	24.559481	0.295752	C
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	C
25%	0.000000	0.000000	0.000000	0.000000	9.000000	1.000000	C
50%	1.000000	0.000000	0.000000	0.000000	29.000000	1.000000	1
75%	1.000000	0.000000	1.000000	1.000000	55.000000	1.000000	2
max	1.000000	1.000000	1.000000	1.000000	72.000000	1.000000	2

In [19]:

```
df.isna().sum()
```

Out[19]:

```
gender                0
SeniorCitizen         0
Partner               0
Dependents            0
tenure                0
PhoneService          0
MultipleLines         0
InternetService       0
OnlineSecurity        0
OnlineBackup          0
DeviceProtection      0
TechSupport           0
StreamingTV           0
StreamingMovies       0
Contract              0
PaperlessBilling      0
PaymentMethod         0
MonthlyCharges        0
TotalCharges          0
Churn                 0
dtype: int64
```

In [20]:

```
x=df.iloc[:,0:-1]
y=df.iloc[:, -1]
```

In [21]:

```
#Feature Selection
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
test=SelectKBest(score_func=chi2)
fi=test.fit(x,y)
fi.scores_
```

Out[21]:

```
array([2.58698618e-01, 1.34351545e+02, 8.24120826e+01, 1.33036443e+02,
       1.62789237e+04, 9.72606249e-02, 9.74692078e+00, 9.82102823e+00,
       5.51611529e+02, 2.30086520e+02, 1.91303140e+02, 5.23303866e+02,
       7.49020319e+00, 8.23539949e+00, 1.11578017e+03, 1.05680863e+02,
       5.84922505e+01, 3.68078770e+03, 4.50421670e+05])
```

In [22]:

```
col=x.columns  
score=pd.DataFrame({'features':col,'score_chi2':fi.scores_})  
score
```

Out[22]:

	features	score_chi2
0	gender	0.258699
1	SeniorCitizen	134.351545
2	Partner	82.412083
3	Dependents	133.036443
4	tenure	16278.923685
5	PhoneService	0.097261
6	MultipleLines	9.746921
7	InternetService	9.821028
8	OnlineSecurity	551.611529
9	OnlineBackup	230.086520
10	DeviceProtection	191.303140
11	TechSupport	523.303866
12	StreamingTV	7.490203
13	StreamingMovies	8.235399
14	Contract	1115.780167
15	PaperlessBilling	105.680863
16	PaymentMethod	58.492250
17	MonthlyCharges	3680.787699
18	TotalCharges	450421.669826

In [23]:

```
score.sort_values(by='score_chi2',ascending=False)
```

Out[23]:

	features	score_chi2
18	TotalCharges	450421.669826
4	tenure	16278.923685
17	MonthlyCharges	3680.787699
14	Contract	1115.780167
8	OnlineSecurity	551.611529
11	TechSupport	523.303866
9	OnlineBackup	230.086520
10	DeviceProtection	191.303140
1	SeniorCitizen	134.351545
3	Dependents	133.036443
15	PaperlessBilling	105.680863
2	Partner	82.412083
16	PaymentMethod	58.492250
7	InternetService	9.821028
6	MultipleLines	9.746921
13	StreamingMovies	8.235399
12	StreamingTV	7.490203
0	gender	0.258699
5	PhoneService	0.097261

In [24]:

```
df.drop(['PhoneService', 'gender', 'StreamingTV', 'StreamingMovies', 'MultipleLines', 'InternetService'], axis=1)
```

Out[24]:

	SeniorCitizen	Partner	Dependents	tenure	OnlineSecurity	OnlineBackup	DeviceProtection
0	0	1	0	1	0	2	
1	0	0	0	34	2	0	
2	0	0	0	2	2	2	
3	0	0	0	45	2	0	
4	0	0	0	2	0	0	
...
7038	0	1	1	24	2	0	
7039	0	1	1	72	0	2	
7040	0	1	1	11	2	0	
7041	1	1	0	4	0	0	
7042	0	0	0	66	2	0	

7043 rows × 14 columns

In [25]:

```
#DATA PREPROCESSING
```

In [26]:

```
#Splitting the dataset into training and testing data
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30, random_state=1)
```

In [27]:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(x_train)
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)
```

In [28]:

```
#Machine Learning Model Prediction and Performance Evaluation
```

In [29]:

#KNN

```

from sklearn.neighbors import KNeighborsClassifier
classifier=KNeighborsClassifier(n_neighbors=9)
classifier.fit(x_train,y_train)
y_pred=classifier.predict(x_test)

```

In [30]:

```

from sklearn.metrics import classification_report,confusion_matrix,accuracy_score,C
print(classification_report(y_test,y_pred))
result=confusion_matrix(y_test,y_pred)
print(result)
print(ConfusionMatrixDisplay.from_predictions(y_test,y_pred))

```


	precision	recall	f1-score	support
0	0.86	0.86	0.86	1585
1	0.57	0.56	0.57	528
accuracy			0.79	2113
macro avg	0.71	0.71	0.71	2113
weighted avg	0.79	0.79	0.79	2113

```

[[1364  221]
 [ 230  298]]

```

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7f3cafadf970>



In [31]:

```
score=accuracy_score(y_test,y_pred)
print(score)
```

0.7865593942262187

In [32]:

```
#Naive Bayes
from sklearn.naive_bayes import GaussianNB
model=GaussianNB()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
```

In [33]:

```

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, C
print(classification_report(y_test, y_pred))
result = confusion_matrix(y_test, y_pred)
print(result)
print(ConfusionMatrixDisplay.from_predictions(y_test, y_pred))


```

	precision	recall	f1-score	support
0	0.91	0.76	0.83	1585
1	0.52	0.77	0.62	528
accuracy			0.77	2113
macro avg	0.72	0.77	0.73	2113
weighted avg	0.81	0.77	0.78	2113

[[1210 375]

[120 408]]

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object
at 0x7f3cabe4cc70>



In [34]:

```

score = accuracy_score(y_test, y_pred)
print(score)

```

0.7657359204921912

In [35]:

```
#SVM
from sklearn.svm import SVC
classifier=SVC()
classifier.fit(x_train,y_train)
y_pred=classifier.predict(x_test)
y_pred
```

Out[35]:


```
array([0, 0, 0, ..., 0, 0, 0])
```

In [36]:

```
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score
print(classification_report(y_test,y_pred))
result=confusion_matrix(y_test,y_pred)
print(result)
print(ConfusionMatrixDisplay.from_predictions(y_test,y_pred))
```

		precision	recall	f1-score	support
	0	0.85	0.91	0.88	1585
	1	0.65	0.51	0.57	528
accuracy				0.81	2113
macro avg		0.75	0.71	0.73	2113
weighted avg		0.80	0.81	0.80	2113

```
[[1443 142]
 [ 259 269]]
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object
at 0x7f3cac868f40>
```



In [37]:

```
score=accuracy_score(y_test,y_pred)
print(score)
```

0.8102224325603408

In [38]:

```
#Decision Tree
from sklearn.tree import DecisionTreeClassifier
dec=DecisionTreeClassifier(criterion='entropy')
dec.fit(x_train,y_train)
y_pred=dec.predict(x_test)
y_pred
```

Out[38]:

array([0, 0, 0, ..., 1, 0, 0])

In [39]:

```

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
print(classification_report(y_test, y_pred))
result = confusion_matrix(y_test, y_pred)
print(result)
print(ConfusionMatrixDisplay.from_predictions(y_test, y_pred))

```

	precision	recall	f1-score	support
0	0.84	0.81	0.82	1585
1	0.48	0.52	0.50	528
accuracy			0.74	2113
macro avg	0.66	0.67	0.66	2113
weighted avg	0.75	0.74	0.74	2113

```


[[1284  301]
 [ 251  277]]

```

```

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object
at 0x7f3cda9ec580>

```



In [40]:

```

score = accuracy_score(y_test, y_pred)
print(score)

```

0.738760056791292

In []:

```

#Comparing the results obtained from each of the above classification algorithms,
#we see that SVM algorithm has the highest accuracy of about 81% for this dataset,
#So we choose SVM classification algorithm for this dataset.

```

